

April 30, 2024

# 1 WEEK 6 FINAL PROJECT - DISTRACTED DRIVER DETECTION USING CNN

Student, University of Colorado Boulder

## 1.1 Gather data, determine the method of data collection and provenance of the data

For the final project, I wanted to work on any object detection/image processing deep learning area. I found a dataset that could help me with this on Kaggle.

While driving on road, it is not uncommon to see drivers who are distracted. They may be texting, scrolling through social media, or in a lively hand-held conversation on their phone.

Quoting from the Kaggle competition: “According to the CDC motor vehicle safety division, one in five car accidents is caused by a distracted driver. Sadly, this translates to 425,000 people injured and 3,000 people killed by distracted driving every year.”

The Kaggle challenge was sponsored by State Farm with the intent to better insure their customers, by testing whether dashboard cameras can automatically detect drivers engaging in distracted behaviors.

The challenge’s outcomes are evaluated using the multi-class logarithmic loss. Each image has been labeled with one true class. For each image, we must submit a set of predicted probabilities (one for every image). The formula is then,

$$\text{logloss} = \frac{-1}{N} \sum_{i=1}^N \sum_{j=1}^M y_{ij} \log(p_{ij}),$$

where

$N$  is the number of images in the test set,

$M$  is the number of image class labels,

$\log$  is the natural logarithm,

$y_{ij}$  is 1 if observation  $i$  belongs to class  $j$  and 0 otherwise,

and  $p_{ij}$  is the predicted probability that observation  $i$  belongs to class  $j$ .

The submitted probabilities for a given image are not required to sum to one because they are rescaled prior to being scored (each row is divided by the row sum). In order to avoid the extremes of the log function, predicted probabilities are replaced with  $\max(\min(p, 1 - 10^{-15}), 10^{-15})$ .

To download the data, we have to register for the competition <https://www.kaggle.com/competitions/state-farm-distracted-driver-detection/overview>

## 1.2 Identify a Deep Learning Problem

For this project, since the data is from Kaggle, I will be trying different CNN models and comparing their performance.

```
[2]: import os, cv2, random, shutil

import pandas as pd
import numpy as np

from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.utils import plot_model

from keras.layers import Conv2D, MaxPooling2D, Dropout, Flatten, Dense
from keras.models import Sequential, load_model
from keras.callbacks import ModelCheckpoint
from keras.utils import to_categorical
from keras.preprocessing import image

from tqdm import tqdm

from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.datasets import load_files
from sklearn.model_selection import train_test_split

%matplotlib inline
```

## 1.3 Exploratory Data Analysis (EDA) - Inspect, Visualize, and Clean the Data

```
[56]: DATA_PATH = "/kaggle"

INPUT_PATH = os.path.join(DATA_PATH, "input/
    state-farm-distracted-driver-detection")
```

```

IMAGES_PATH = os.path.join(INPUT_PATH, "imgs")
TEST_DIR = os.path.join(IMAGES_PATH, "test")
TRAIN_DIR = os.path.join(IMAGES_PATH, "train")

BASE_MODEL_PATH = os.path.join(DATA_PATH, "working", "model")

MODEL_PATH = os.path.join(BASE_MODEL_PATH, "self_trained")
OUTPUT_CSV_FILES_DIR = os.path.join(DATA_PATH, "working", "output_csv_files")
INPUT_CSV_FILES_DIR = os.path.join(DATA_PATH, "working", "input_csv_files")

if not os.path.exists(INPUT_CSV_FILES_DIR):
    os.makedirs(INPUT_CSV_FILES_DIR)
    print(f"Input CSV path created {INPUT_CSV_FILES_DIR}")

if not os.path.exists(OUTPUT_CSV_FILES_DIR):
    os.makedirs(OUTPUT_CSV_FILES_DIR)
    print(f"Output CSV path created {OUTPUT_CSV_FILES_DIR}")

if not os.path.exists(MODEL_PATH):
    os.makedirs(MODEL_PATH)
    print(f"Model path created {MODEL_PATH}")

```

```

[20]: train_image = []
      image_label = []

      img_dim = 64

      driver_details = pd.read_csv(os.path.join(INPUT_PATH, "driver_imgs_list.csv"),
      ↪na_values='na')
      driver_details.head(5)

      for i in range(10):
          print(f'traversing C{i}')
          imgs = os.listdir(os.path.join(TRAIN_DIR, "c"+str(i)))
          for j in range(len(imgs)):
              img_name = os.path.join(TRAIN_DIR, "c"+str(i), imgs[j])
              img = cv2.imread(img_name)
              img = img[50:,120:-50]
              img = cv2.resize(img,(img_dim,img_dim))
              label = i
              driver = driver_details[driver_details['img'] == imgs[j]]['subject'].
              ↪values[0]
              train_image.append([img,label,driver])
              image_label.append(i)

```

```

traversing C0
traversing C1
traversing C2

```

```
traversing C3
traversing C4
traversing C5
traversing C6
traversing C7
traversing C8
traversing C9
```

```
[65]: driver_details.head(5)
```

```
[65]:  subject  classname      img
      0    p002      c0  img_44733.jpg
      1    p002      c0  img_72999.jpg
      2    p002      c0  img_25094.jpg
      3    p002      c0  img_69092.jpg
      4    p002      c0  img_92629.jpg
```

```
[66]: driver_details.describe()
```

```
[66]:  subject  classname      img
count    22424      22424      22424
unique      26         10      22424
top        p021      c0  img_9684.jpg
freq       1237      2489         1
```

```
[68]: driver_details.isna().sum()
```

```
[68]: subject      0
      classname    0
      img          0
      dtype: int64
```

All images seem to be good. We have no null values in csv. So, there is nothing to be cleaned

```
[21]: class_value_counts = dict()
      for img in train_image:
          class_type = img[1]
          class_value_counts['c'+str(class_type)] = class_value_counts.
          ↳get('c'+str(class_type), 0)+1

      print(class_value_counts)
```

```
{'c0': 2489, 'c1': 2267, 'c2': 2317, 'c3': 2346, 'c4': 2326, 'c5': 2312, 'c6':
2325, 'c7': 2002, 'c8': 1911, 'c9': 2129}
```


Let us look at the number of images in each of the different classes

```
[22]: fig = plt.figure(figsize=(18,12))
ay = fig.add_subplot(211)
labels = class_value_counts.keys()
x = range(len(labels))
y = class_value_counts.values()
plt.xticks(x, labels, size=10)
plt.yticks(size=10)

ay.bar(x, y, color="blue")

plt.title('Count of Images by Classification',size=20)
plt.xlabel('Classification', size=15)
plt.ylabel('Count', size=15)

plt.show()
```




Let us look at one image from each of the different classes

```
[60]: fig, ax = plt.subplots(1, 10, figsize = (50,50 ))
for i in range(10):
    for file in os.listdir(os.path.join(TRAIN_DIR, 'c'+str(i))):
        img = cv2.imread(os.path.join(TRAIN_DIR, 'c'+str(i),file))
        img = img[50:,120:-50]
        img = cv2.resize(img,(224,224))

        ax[i].imshow(img,cmap = 'gray')
    break
plt.show
```

```
[60]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
[74]: # Find the frequency of images per driver
driver_ids = dict()
for dr in driver_details['subject']:
    driver_ids[dr] = driver_ids.get(dr, 0)+1


d = list()
c = list()
for k,v in driver_ids.items():
    d.append(k)
    c.append(v)

fig = plt.figure(figsize=(18,12))
ay = fig.add_subplot(211)
labels = d
x = range(len(labels))
y = c
plt.xticks(x, labels, size=10)
plt.yticks(size=10)

ay.bar(x, y, color="blue")

plt.title('Count of Images by Driver',size=20)
plt.xlabel('Driver', size=15)
plt.ylabel('Count', size=15)

plt.show()
```



## 1.4 Model Building

I will split the data into 80:20 for training and validation. The CNN model I will start off with is a multi layer one with Conv2D with 64+128+256+512 filters and relu activation, followed by max

pooling. After flattening the data and adding dropouts, finally I will add a softmax layer with 10 classes.

```
[32]: X = list()
y = list()
for features, labels, drivers in train_image:
    X.append(features)
    y.append(labels)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)
print (len(X_train), len(X_test))
print (len(y_train), len(y_test))
X_train = np.array(X_train).reshape(-1, img_dim, img_dim, 3)
X_test = np.array(X_test).reshape(-1, img_dim, img_dim, 3)
print(f'{X_train.shape} {X_test.shape}')
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
print(f'{y_train.shape} {y_test.shape}')
```

17939 4485

17939 4485

(17939, 64, 64, 3) (4485, 64, 64, 3)

(17939, 10) (4485, 10)

```
[33]: def get_model():
    model = Sequential()
    # 64 conv2d filters with relu
    model.add(Conv2D(filters=64, kernel_size=2, padding='same',
    ↪activation='relu', input_shape=(64,64,3),
    ↪kernel_initializer='glorot_normal'))
    model.add(MaxPooling2D(pool_size=2))

    # 128 conv2d filters with relu
    model.add(Conv2D(filters=128, kernel_size=2, padding='same',
    ↪activation='relu', kernel_initializer='glorot_normal'))
    model.add(MaxPooling2D(pool_size=2))

    # 256 conv2d filters with relu
    model.add(Conv2D(filters=256, kernel_size=2, padding='same',
    ↪activation='relu', kernel_initializer='glorot_normal'))
    model.add(MaxPooling2D(pool_size=2))

    # 512 conv2d filters with relu
    model.add(Conv2D(filters=512, kernel_size=2, padding='same',
    ↪activation='relu', kernel_initializer='glorot_normal'))
    model.add(MaxPooling2D(pool_size=2))
```

```

        model.add(Dropout(0.5))
        model.add(Flatten())
        model.add(Dense(500, activation='relu', kernel_initializer='glorot_normal'))
        model.add(Dropout(0.5))
        model.add(Dense(10, activation='softmax',
↪kernel_initializer='glorot_normal'))

    return model

model = get_model()
model.summary()

```

/opt/conda/lib/python3.10/site-packages/keras/src/layers/convolutional/base\_conv.py:99: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 64, 64, 64)	832
max_pooling2d_4 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_5 (Conv2D)	(None, 32, 32, 128)	32,896
max_pooling2d_5 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_6 (Conv2D)	(None, 16, 16, 256)	131,328
max_pooling2d_6 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_7 (Conv2D)	(None, 8, 8, 512)	524,800
max_pooling2d_7 (MaxPooling2D)	(None, 4, 4, 512)	0
dropout_2 (Dropout)	(None, 4, 4, 512)	0
flatten_1 (Flatten)	(None, 8192)	0
dense_2 (Dense)	(None, 500)	4,096,500
dropout_3 (Dropout)	(None, 500)	0



dense\_3 (Dense) (None, 10) 5,010

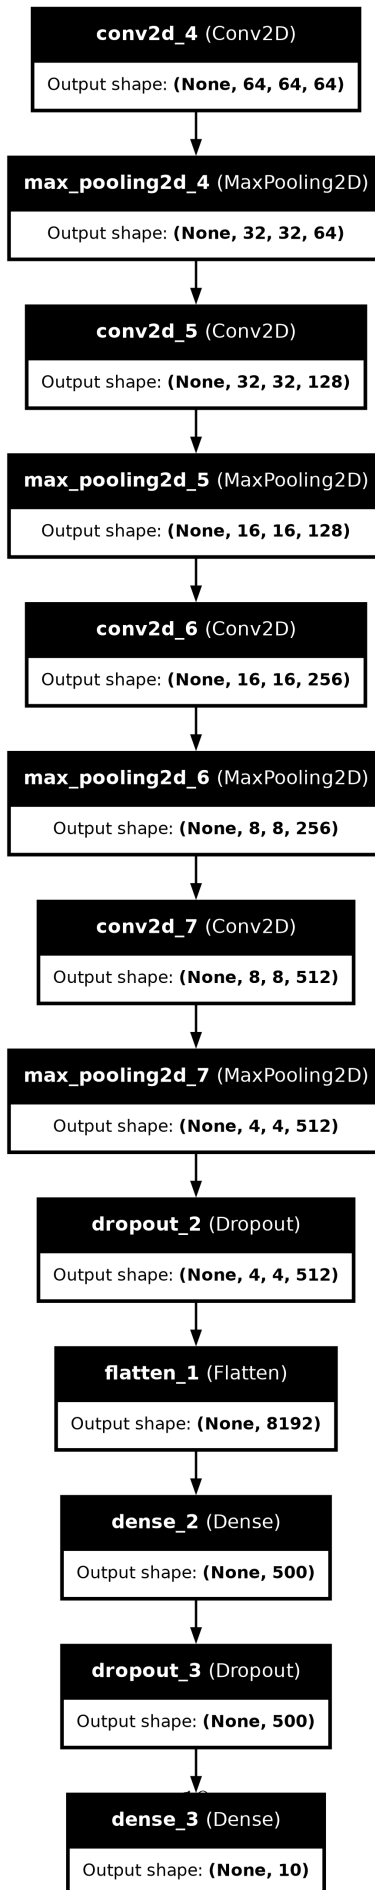
Total params: 4,791,366 (18.28 MB)

Trainable params: 4,791,366 (18.28 MB)

Non-trainable params: 0 (0.00 B)

```
[34]: plot_model(model, to_file=os.path.join(MODEL_PATH, "model_distracted_driver.  
      ↳png"), show_shapes=True, show_layer_names=True)
```

[34]:



```
[35]: model.compile(optimizer='rmsprop', loss='categorical_crossentropy',  
    ↪ metrics=['accuracy'])  
  
    filepath = os.path.join(MODEL_PATH, "first-{epoch:02d}-{val_accuracy:.2f}.keras")  
  
    img_data_gen = ImageDataGenerator(height_shift_range=0.5, width_shift_range=0.  
    ↪ 5, zoom_range=0.5, rotation_range=30)  
  
    data_generator = img_data_gen.flow(X_train, y_train, batch_size=64)  
  
    checkpoint = ModelCheckpoint(filepath, monitor='val_accuracy', verbose=1,  
    ↪ save_best_only=True, mode='max')  
    callbacks_list = [checkpoint]  
    epochs = 20  
  
    model_history = model.fit(data_generator, validation_data=(X_test, y_test),  
    ↪ epochs=epochs, batch_size=40, shuffle=True, callbacks=callbacks_list)
```

Epoch 1/20

```
/opt/conda/lib/python3.10/site-  
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:120:  
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in  
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,  
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be  
ignored.
```

```
    self._warn_if_super_not_called()
```

```
3/281          17s 62ms/step - accuracy:  
0.0703 - loss: 94.4282
```

```
WARNING: All log messages before absl::InitializeLog() is called are written to  
STDERR
```

```
I0000 00:00:1714031376.956743      87 device_compiler.h:186] Compiled cluster  
using XLA! This line is logged at most once for the lifetime of the process.  
W0000 00:00:1714031376.973145      87 graph_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update
```

```
118/281          20s 125ms/step -  
accuracy: 0.1038 - loss: 14.2558
```

```
W0000 00:00:1714031391.547976      90 graph_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update
```

```
280/281          0s 102ms/step -  
accuracy: 0.1086 - loss: 8.3521
```

```
W0000 00:00:1714031407.352311      89 graph_launch.cc:671] Fallback to op-by-op
```

mode because memset node breaks graph update

Epoch 1: val\_accuracy improved from -inf to 0.22965, saving model to  
/kaggle/working/model/self\_trained/first-01-0.23.keras

281/281 42s 114ms/step -

accuracy: 0.1087 - loss: 8.3175 - val\_accuracy: 0.2297 - val\_loss: 2.1536

W0000 00:00:1714031408.718173 88 graph\_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update

Epoch 2/20

281/281 0s 87ms/step -

accuracy: 0.1476 - loss: 2.2606

Epoch 2: val\_accuracy did not improve from 0.22965

281/281 25s 88ms/step -

accuracy: 0.1477 - loss: 2.2605 - val\_accuracy: 0.2161 - val\_loss: 2.1655

Epoch 3/20

280/281 0s 85ms/step -

accuracy: 0.1801 - loss: 2.1764

Epoch 3: val\_accuracy improved from 0.22965 to 0.30814, saving model to  
/kaggle/working/model/self\_trained/first-03-0.31.keras

281/281 25s 87ms/step -

accuracy: 0.1802 - loss: 2.1762 - val\_accuracy: 0.3081 - val\_loss: 1.8722

Epoch 4/20

281/281 0s 85ms/step -

accuracy: 0.2192 - loss: 2.0706

Epoch 4: val\_accuracy improved from 0.30814 to 0.41427, saving model to  
/kaggle/working/model/self\_trained/first-04-0.41.keras

281/281 25s 87ms/step -

accuracy: 0.2193 - loss: 2.0706 - val\_accuracy: 0.4143 - val\_loss: 1.5288

Epoch 5/20

280/281 0s 86ms/step -

accuracy: 0.2676 - loss: 1.9692

Epoch 5: val\_accuracy improved from 0.41427 to 0.42876, saving model to  
/kaggle/working/model/self\_trained/first-05-0.43.keras

281/281 25s 88ms/step -

accuracy: 0.2676 - loss: 1.9692 - val\_accuracy: 0.4288 - val\_loss: 1.4417

Epoch 6/20

281/281 0s 84ms/step -

accuracy: 0.2957 - loss: 1.9029

Epoch 6: val\_accuracy improved from 0.42876 to 0.52263, saving model to  
/kaggle/working/model/self\_trained/first-06-0.52.keras

281/281 25s 86ms/step -

accuracy: 0.2957 - loss: 1.9028 - val\_accuracy: 0.5226 - val\_loss: 1.3904

Epoch 7/20

280/281 0s 85ms/step -

accuracy: 0.3237 - loss: 1.8160

Epoch 7: val\_accuracy did not improve from 0.52263

281/281                    25s 87ms/step -  
accuracy: 0.3238 - loss: 1.8160 - val\_accuracy: 0.5119 - val\_loss: 1.3034  
Epoch 8/20  
280/281                    0s 84ms/step -  
accuracy: 0.3498 - loss: 1.7702  
Epoch 8: val\_accuracy improved from 0.52263 to 0.58016, saving model to  
/kaggle/working/model/self\_trained/first-08-0.58.keras  
281/281                    25s 86ms/step -  
accuracy: 0.3498 - loss: 1.7702 - val\_accuracy: 0.5802 - val\_loss: 1.2302  
Epoch 9/20  
281/281                    0s 85ms/step -  
accuracy: 0.3596 - loss: 1.7301  
Epoch 9: val\_accuracy did not improve from 0.58016  
281/281                    25s 86ms/step -  
accuracy: 0.3596 - loss: 1.7301 - val\_accuracy: 0.5710 - val\_loss: 1.2306  
Epoch 10/20  
280/281                    0s 86ms/step -  
accuracy: 0.3777 - loss: 1.7069  
Epoch 10: val\_accuracy improved from 0.58016 to 0.64169, saving model to  
/kaggle/working/model/self\_trained/first-10-0.64.keras  
281/281                    25s 88ms/step -  
accuracy: 0.3778 - loss: 1.7068 - val\_accuracy: 0.6417 - val\_loss: 1.0969  
Epoch 11/20  
280/281                    0s 84ms/step -  
accuracy: 0.3860 - loss: 1.6998  
Epoch 11: val\_accuracy improved from 0.64169 to 0.67826, saving model to  
/kaggle/working/model/self\_trained/first-11-0.68.keras  
281/281                    25s 86ms/step -  
accuracy: 0.3861 - loss: 1.6998 - val\_accuracy: 0.6783 - val\_loss: 0.8630  
Epoch 12/20  
280/281                    0s 84ms/step -  
accuracy: 0.3956 - loss: 1.6780  
Epoch 12: val\_accuracy improved from 0.67826 to 0.69387, saving model to  
/kaggle/working/model/self\_trained/first-12-0.69.keras  
281/281                    25s 86ms/step -  
accuracy: 0.3957 - loss: 1.6779 - val\_accuracy: 0.6939 - val\_loss: 0.8647  
Epoch 13/20  
280/281                    0s 86ms/step -  
accuracy: 0.4205 - loss: 1.6266  
Epoch 13: val\_accuracy did not improve from 0.69387  
281/281                    25s 88ms/step -  
accuracy: 0.4205 - loss: 1.6266 - val\_accuracy: 0.5960 - val\_loss: 1.3777  
Epoch 14/20  
281/281                    0s 86ms/step -  
accuracy: 0.4155 - loss: 1.6389  
Epoch 14: val\_accuracy did not improve from 0.69387  
281/281                    25s 87ms/step -  
accuracy: 0.4155 - loss: 1.6389 - val\_accuracy: 0.6736 - val\_loss: 0.9780

```

Epoch 15/20
280/281          0s 85ms/step -
accuracy: 0.4434 - loss: 1.5828
Epoch 15: val_accuracy improved from 0.69387 to 0.72441, saving model to
/kaggle/working/model/self_trained/first-15-0.72.keras
281/281          25s 87ms/step -
accuracy: 0.4433 - loss: 1.5828 - val_accuracy: 0.7244 - val_loss: 0.7616
Epoch 16/20
281/281          0s 85ms/step -
accuracy: 0.4432 - loss: 1.5894
Epoch 16: val_accuracy improved from 0.72441 to 0.72843, saving model to
/kaggle/working/model/self_trained/first-16-0.73.keras
281/281          25s 87ms/step -
accuracy: 0.4432 - loss: 1.5894 - val_accuracy: 0.7284 - val_loss: 0.8195
Epoch 17/20
280/281          0s 87ms/step -
accuracy: 0.4434 - loss: 1.5874
Epoch 17: val_accuracy improved from 0.72843 to 0.73043, saving model to
/kaggle/working/model/self_trained/first-17-0.73.keras
281/281          25s 88ms/step -
accuracy: 0.4434 - loss: 1.5874 - val_accuracy: 0.7304 - val_loss: 1.0802
Epoch 18/20
281/281          0s 84ms/step -
accuracy: 0.4512 - loss: 1.6058
Epoch 18: val_accuracy improved from 0.73043 to 0.79465, saving model to
/kaggle/working/model/self_trained/first-18-0.79.keras
281/281          25s 86ms/step -
accuracy: 0.4512 - loss: 1.6058 - val_accuracy: 0.7946 - val_loss: 0.7764
Epoch 19/20
281/281          0s 87ms/step -
accuracy: 0.4460 - loss: 1.5926
Epoch 19: val_accuracy did not improve from 0.79465
281/281          25s 88ms/step -
accuracy: 0.4461 - loss: 1.5925 - val_accuracy: 0.7672 - val_loss: 0.7132
Epoch 20/20
280/281          0s 86ms/step -
accuracy: 0.4535 - loss: 1.5571
Epoch 20: val_accuracy did not improve from 0.79465
281/281          25s 87ms/step -
accuracy: 0.4536 - loss: 1.5571 - val_accuracy: 0.6974 - val_loss: 1.4080

```

```

[38]: def show_loss_and_accuracy_plots(model_history):
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 12))
    ax1.plot(model_history.history['loss'], color='b', label="Training Loss")
    ax1.plot(model_history.history['val_loss'], color='r', label="Validation_
↳ Loss")
    ax1.set_xticks(np.arange(1, 25, 1))

```


```

ax2.plot(model_history.history['accuracy'], color='b', label="Training_
↪Accuracy")
ax2.plot(model_history.history['val_accuracy'], color='r',label="Validation_
↪Accuracy")
ax2.set_xticks(np.arange(1, 25, 1))

legend = plt.legend(loc='best', shadow=True)
plt.tight_layout()
plt.show()

show_loss_and_accuracy_plots(model_history)

```



```
[39]: def print_confusion_matrix(confusion_matrix, class_names, figsize=(10,7),
    ↪fontsize=14):
    df_cm = pd.DataFrame(
        confusion_matrix, index=class_names, columns=class_names,
    )
    fig = plt.figure(figsize=figsize)
    try:
        heatmap = sns.heatmap(df_cm, annot=True, fmt="d")
    except ValueError:
        raise ValueError("Confusion matrix values must be integers.")
    heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0,
    ↪ha='right', fontsize=fontsize)
    heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=45,
    ↪ha='right', fontsize=fontsize)
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    fig.savefig(os.path.join(MODEL_PATH, "confusion_matrix.png"))
    return fig

def print_heatmap(n_labels, n_predictions, class_names):
    labels = n_labels
    predictions = n_predictions

    matrix = confusion_matrix(labels.argmax(axis=1), predictions.argmax(axis=1))
    row_sum = np.sum(matrix, axis = 1)
    w, h = matrix.shape
    c_m = np.zeros((w, h))
    for i in range(h):
        c_m[i] = matrix[i] * 100 / row_sum[i]
    c = c_m.astype(dtype = np.uint8)
    heatmap = print_confusion_matrix(c, class_names, figsize=(12,8),
    ↪fontsize=16)

class_names = ['c0', 'c1', 'c2', 'c3', 'c4', 'c5', 'c6', 'c7', 'c8', 'c9']
y_pred = model.predict(X_test)

print_heatmap(y_test, y_pred, class_names)
```


66/141                      0s 2ms/step

W0000 00:00:1714032332.210508            90 graph\_launch.cc:671] Fallback to op-by-op mode because memset node breaks graph update

141/141                      2s 5ms/step

W0000 00:00:1714032332.911818            89 graph\_launch.cc:671] Fallback to op-by-op mode because memset node breaks graph update





```
[41]: y_pred_class = np.argmax(y_pred, axis=1)

y_test = np.argmax(y_test, axis=1)

accuracy = accuracy_score(y_test, y_pred_class)
print('Accuracy: %f' % accuracy)

precision = precision_score(y_test, y_pred_class, average='weighted')
print('Precision: %f' % precision)

recall = recall_score(y_test, y_pred_class, average='weighted')
print('Recall: %f' % recall)

f1 = f1_score(y_test, y_pred_class, average='weighted')
print('F1 score: %f' % f1)
```

```
Accuracy: 0.697436
Precision: 0.744432
Recall: 0.697436
F1 score: 0.683957
```

```
[42]: #load checkpoint with best val_accuracy
BEST_MODEL = os.path.join(MODEL_PATH, "first-18-0.79.keras")
model = load_model(BEST_MODEL)
model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 64, 64, 64)	832
max_pooling2d_4 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_5 (Conv2D)	(None, 32, 32, 128)	32,896
max_pooling2d_5 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_6 (Conv2D)	(None, 16, 16, 256)	131,328
max_pooling2d_6 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_7 (Conv2D)	(None, 8, 8, 512)	524,800
max_pooling2d_7 (MaxPooling2D)	(None, 4, 4, 512)	0
dropout_2 (Dropout)	(None, 4, 4, 512)	0
flatten_1 (Flatten)	(None, 8192)	0
dense_2 (Dense)	(None, 500)	4,096,500
dropout_3 (Dropout)	(None, 500)	0
dense_3 (Dense)	(None, 10)	5,010

Total params: 9,582,734 (36.56 MB)

Trainable params: 4,791,366 (18.28 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 4,791,368 (18.28 MB)

```
[43]: files = os.listdir(TEST_DIR)
test_images = []
for i in range(len(files)):
    img = cv2.imread(os.path.join(TEST_DIR, files[i]))
    img = img[50:,120:-50]
    img = cv2.resize(img,(64,64))
    test_images.append(img)

test_images = np.array(test_images).reshape(-1,64,64,3)
test_image.shape
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[43], line 10
      7     test_images.append(img)
      9 test_images = np.array(test_images).reshape(-1,64,64,3)
----> 10 test_image.shape

NameError: name 'test_image' is not defined
```

```
[44]: test_images.shape
```

```
[44]: (79726, 64, 64, 3)
```

```
[47]: def predict_and_write_results(model, test_images, files, out_file_name):
    df = pd.DataFrame(columns=['img', 'c0', 'c1', 'c2', 'c3', 'c4', 'c5', 'c6', 'c7', 'c8', 'c9'])
    ypred_test = model.predict(test_images, verbose=1)
    for i in range(len(files)):
        new_row = {
            "img":files[i],
            'c0':ypred_test[i][0],
            'c1':ypred_test[i][1],
            'c2':ypred_test[i][2],
            'c3':ypred_test[i][3],
            'c4':ypred_test[i][4],
            'c5':ypred_test[i][5],
            'c6':ypred_test[i][6],
            'c7':ypred_test[i][7],
            'c8':ypred_test[i][8],
            'c9':ypred_test[i][9]
        }
        df = pd.concat([df, pd.DataFrame([new_row])], ignore_index=True)
    df.to_csv(os.path.join(OUTPUT_CSV_FILES_DIR,out_file_name),index=False)
```




```
predict_and_write_results(model, test_images, files,
↪ 'kaggle_self_trained_first_model.csv')
```

2492/2492

6s 2ms/step

/tmp/ipykernel\_34/950582733.py:18: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```
df = pd.concat([df, pd.DataFrame([new_row])], ignore_index=True)
```

Submission and Description	Private Score 	Public Score 	Selected
 <b>kaggle_self_trained_first_model.csv</b> <small>Complete (after deadline) · now</small>	<b>2.24537</b>	<b>2.46848</b>	<input type="checkbox"/>

## 1.5 Tuning training process

This score is good. However, when the images are split randomly, different images of the same person will end up in training and validation and there appears to be a leak in the model. Let us try if the same model will work better if we split our training data according to driver instead of random split

```
[48]: driv_selected = ['p050', 'p015', 'p022', 'p056']
      ## Split the data into training and validation
```

```
X_train= []
y_train = []
X_test = []
y_test = []
D_train = []
D_test = []

for features,labels,drivers in train_image:
    if drivers in driv_selected:
        X_test.append(features)
        y_test.append(labels)
        D_test.append(drivers)

    else:
        X_train.append(features)
        y_train.append(labels)
        D_train.append(drivers)

print (len(X_train),len(X_test))
print (len(y_train),len(y_test))
```

```
18732 3692
18732 3692
```

```
[49]: X_train = np.array(X_train).reshape(-1,img_dim,img_dim,3)
X_test = np.array(X_test).reshape(-1,img_dim,img_dim,3)
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

print (X_train.shape)
```

```
(18732, 64, 64, 3)
```

```
[50]: model = get_model()
model.summary()
```

```
/opt/conda/lib/python3.10/site-
packages/keras/src/layers/convolutional/base_conv.py:99: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
```

```
super().__init__()
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 64, 64, 64)	832
max_pooling2d_8 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_9 (Conv2D)	(None, 32, 32, 128)	32,896
max_pooling2d_9 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_10 (Conv2D)	(None, 16, 16, 256)	131,328
max_pooling2d_10 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_11 (Conv2D)	(None, 8, 8, 512)	524,800
max_pooling2d_11 (MaxPooling2D)	(None, 4, 4, 512)	0
dropout_4 (Dropout)	(None, 4, 4, 512)	0
flatten_2 (Flatten)	(None, 8192)	0
dense_4 (Dense)	(None, 500)	4,096,500

dropout_5 (Dropout)	(None, 500)	0
dense_5 (Dense)	(None, 10)	5,010

Total params: 4,791,366 (18.28 MB)

Trainable params: 4,791,366 (18.28 MB)

Non-trainable params: 0 (0.00 B)

```
[51]: model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
    ↪ metrics=['accuracy'])

filepath = os.path.join(MODEL_PATH, "second-{epoch:02d}-{val_accuracy:.2f}.
    ↪ keras")

checkpoint = ModelCheckpoint(filepath, monitor='val_accuracy', verbose=1,
    ↪ save_best_only=True, mode='max')
callbacks_list = [checkpoint]

img_data_gen = ImageDataGenerator(height_shift_range=0.5, width_shift_range = 0.
    ↪ 5, zoom_range = 0.5, rotation_range=30)

data_generator = img_data_gen.flow(X_train, y_train, batch_size=64)

history = model.fit(data_generator, callbacks=callbacks_list, epochs=20,
    ↪ verbose=1, validation_data=(X_test, y_test), batch_size=40, shuffle=True)
```

Epoch 1/20

```
/opt/conda/lib/python3.10/site-
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:120:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
  self._warn_if_super_not_called()

3/293          17s 62ms/step - accuracy:
0.0972 - loss: 76.3583

W0000 00:00:1714037371.592966      87 graph_launch.cc:671] Fallback to op-by-op
mode because memset node breaks graph update
```

124/293                    22s 134ms/step -  
accuracy: 0.0970 - loss: 11.9435

W0000 00:00:1714037387.950956            90 graph\_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update

293/293                    0s 106ms/step -  
accuracy: 0.1049 - loss: 7.2001

W0000 00:00:1714037403.532477            89 graph\_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update

Epoch 1: val\_accuracy improved from -inf to 0.18743, saving model to  
/kaggle/working/model/self\_trained/second-01-0.19.keras

293/293                    39s 115ms/step -  
accuracy: 0.1049 - loss: 7.1867 - val\_accuracy: 0.1874 - val\_loss: 2.1930

Epoch 2/20

W0000 00:00:1714037404.999101            87 graph\_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update

292/293                    0s 86ms/step -  
accuracy: 0.1452 - loss: 2.2601

Epoch 2: val\_accuracy did not improve from 0.18743

293/293                    26s 87ms/step -  
accuracy: 0.1452 - loss: 2.2600 - val\_accuracy: 0.1603 - val\_loss: 2.6819

Epoch 3/20

293/293                    0s 86ms/step -  
accuracy: 0.1961 - loss: 2.1519

Epoch 3: val\_accuracy improved from 0.18743 to 0.32990, saving model to  
/kaggle/working/model/self\_trained/second-03-0.33.keras

293/293                    26s 88ms/step -  
accuracy: 0.1961 - loss: 2.1518 - val\_accuracy: 0.3299 - val\_loss: 1.8668

Epoch 4/20

292/293                    0s 87ms/step -  
accuracy: 0.2519 - loss: 2.0038

Epoch 4: val\_accuracy improved from 0.32990 to 0.36809, saving model to  
/kaggle/working/model/self\_trained/second-04-0.37.keras

293/293                    27s 89ms/step -  
accuracy: 0.2520 - loss: 2.0037 - val\_accuracy: 0.3681 - val\_loss: 1.8019

Epoch 5/20

292/293                    0s 86ms/step -  
accuracy: 0.2955 - loss: 1.9057

Epoch 5: val\_accuracy improved from 0.36809 to 0.37405, saving model to  
/kaggle/working/model/self\_trained/second-05-0.37.keras

293/293                    26s 87ms/step -  
accuracy: 0.2955 - loss: 1.9055 - val\_accuracy: 0.3741 - val\_loss: 1.8558

Epoch 6/20

292/293                    0s 86ms/step -  
accuracy: 0.3329 - loss: 1.8245

Epoch 6: val\_accuracy did not improve from 0.37405  
293/293 26s 87ms/step -  
accuracy: 0.3329 - loss: 1.8244 - val\_accuracy: 0.3730 - val\_loss: 2.3810  
Epoch 7/20  
292/293 0s 87ms/step -  
accuracy: 0.3574 - loss: 1.7612  
Epoch 7: val\_accuracy did not improve from 0.37405  
293/293 26s 88ms/step -  
accuracy: 0.3575 - loss: 1.7611 - val\_accuracy: 0.3421 - val\_loss: 2.8826  
Epoch 8/20  
292/293 0s 86ms/step -  
accuracy: 0.3840 - loss: 1.6995  
Epoch 8: val\_accuracy improved from 0.37405 to 0.44475, saving model to  
/kaggle/working/model/self\_trained/second-08-0.44.keras  
293/293 26s 87ms/step -  
accuracy: 0.3841 - loss: 1.6993 - val\_accuracy: 0.4447 - val\_loss: 2.7669  
Epoch 9/20  
292/293 0s 86ms/step -  
accuracy: 0.4049 - loss: 1.6577  
Epoch 9: val\_accuracy did not improve from 0.44475  
293/293 26s 87ms/step -  
accuracy: 0.4049 - loss: 1.6576 - val\_accuracy: 0.4309 - val\_loss: 2.2379  
Epoch 10/20  
292/293 0s 86ms/step -  
accuracy: 0.4282 - loss: 1.5863  
Epoch 10: val\_accuracy did not improve from 0.44475  
293/293 26s 87ms/step -  
accuracy: 0.4282 - loss: 1.5862 - val\_accuracy: 0.3922 - val\_loss: 2.1245  
Epoch 11/20  
292/293 0s 85ms/step -  
accuracy: 0.4326 - loss: 1.5790  
Epoch 11: val\_accuracy improved from 0.44475 to 0.48619, saving model to  
/kaggle/working/model/self\_trained/second-11-0.49.keras  
293/293 26s 86ms/step -  
accuracy: 0.4327 - loss: 1.5788 - val\_accuracy: 0.4862 - val\_loss: 1.9426  
Epoch 12/20  
293/293 0s 86ms/step -  
accuracy: 0.4553 - loss: 1.5594  
Epoch 12: val\_accuracy did not improve from 0.48619  
293/293 26s 87ms/step -  
accuracy: 0.4553 - loss: 1.5593 - val\_accuracy: 0.4388 - val\_loss: 2.7675  
Epoch 13/20  
293/293 0s 86ms/step -  
accuracy: 0.4757 - loss: 1.4924  
Epoch 13: val\_accuracy did not improve from 0.48619  
293/293 26s 87ms/step -  
accuracy: 0.4757 - loss: 1.4924 - val\_accuracy: 0.4832 - val\_loss: 2.3321  
Epoch 14/20




```

293/293          0s 86ms/step -
accuracy: 0.4754 - loss: 1.4867
Epoch 14: val_accuracy improved from 0.48619 to 0.55796, saving model to
/kaggle/working/model/self_trained/second-14-0.56.keras
293/293          26s 87ms/step -
accuracy: 0.4754 - loss: 1.4867 - val_accuracy: 0.5580 - val_loss: 2.1365
Epoch 15/20
292/293          0s 85ms/step -
accuracy: 0.4922 - loss: 1.4839
Epoch 15: val_accuracy did not improve from 0.55796
293/293          26s 86ms/step -
accuracy: 0.4922 - loss: 1.4838 - val_accuracy: 0.4783 - val_loss: 5.2060
Epoch 16/20
293/293          0s 86ms/step -
accuracy: 0.4852 - loss: 1.4737
Epoch 16: val_accuracy did not improve from 0.55796
293/293          26s 87ms/step -
accuracy: 0.4852 - loss: 1.4737 - val_accuracy: 0.5447 - val_loss: 1.8665
Epoch 17/20
293/293          0s 86ms/step -
accuracy: 0.4961 - loss: 1.4276
Epoch 17: val_accuracy did not improve from 0.55796
293/293          26s 87ms/step -
accuracy: 0.4962 - loss: 1.4276 - val_accuracy: 0.5087 - val_loss: 2.8459
Epoch 18/20
292/293          0s 86ms/step -
accuracy: 0.5116 - loss: 1.4315
Epoch 18: val_accuracy improved from 0.55796 to 0.56582, saving model to
/kaggle/working/model/self_trained/second-18-0.57.keras
293/293          26s 87ms/step -
accuracy: 0.5116 - loss: 1.4317 - val_accuracy: 0.5658 - val_loss: 1.5440
Epoch 19/20
291/293          0s 87ms/step -
accuracy: 0.4902 - loss: 1.5011
Epoch 19: val_accuracy did not improve from 0.56582
293/293          26s 87ms/step -
accuracy: 0.4902 - loss: 1.5010 - val_accuracy: 0.5309 - val_loss: 3.0278
Epoch 20/20
293/293          0s 85ms/step -
accuracy: 0.4959 - loss: 1.4667
Epoch 20: val_accuracy improved from 0.56582 to 0.58586, saving model to
/kaggle/working/model/self_trained/second-20-0.59.keras
293/293          26s 87ms/step -
accuracy: 0.4960 - loss: 1.4666 - val_accuracy: 0.5859 - val_loss: 1.5672

```

```
[53]: show_loss_and_accuracy_plots(history)
```



```
[52]: #load checkpoint with best val_accuracy
BEST_MODEL = os.path.join(MODEL_PATH, "second-20-0.59.keras")
model = load_model(BEST_MODEL)
model.summary()
predict_and_write_results(model, test_images, files,
    ↪ 'kaggle_self_trained_second_model.csv')
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 64, 64, 64)	832

max_pooling2d_8 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_9 (Conv2D)	(None, 32, 32, 128)	32,896
max_pooling2d_9 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_10 (Conv2D)	(None, 16, 16, 256)	131,328
max_pooling2d_10 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_11 (Conv2D)	(None, 8, 8, 512)	524,800
max_pooling2d_11 (MaxPooling2D)	(None, 4, 4, 512)	0
dropout_4 (Dropout)	(None, 4, 4, 512)	0
flatten_2 (Flatten)	(None, 8192)	0
dense_4 (Dense)	(None, 500)	4,096,500
dropout_5 (Dropout)	(None, 500)	0
dense_5 (Dense)	(None, 10)	5,010

Total params: 9,582,734 (36.56 MB)

Trainable params: 4,791,366 (18.28 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 4,791,368 (18.28 MB)

42/2492                      9s 4ms/step

W0000 00:00:1714038070.214186            87 graph\_launch.cc:671] Fallback to op-by-op mode because memset node breaks graph update

2492/2492                      7s 3ms/step

W0000 00:00:1714038076.708742            90 graph\_launch.cc:671] Fallback to op-by-op mode because memset node breaks graph update

/tmp/ipykernel\_34/950582733.py:18: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the

concat operation.

```
df = pd.concat([df, pd.DataFrame([new_row])], ignore_index=True)
```


Submission and Description		Private Score	Public Score	Selected
	kaggle_self_trained_second_model.csv Complete (after deadline) - now	1.62022	1.71969	<input type="checkbox"/>

```
[54]: y_pred = model.predict(X_test)

print_heatmap(y_test,y_pred,class_names)
```

116/116 1s 6ms/step

W0000 00:00:1714041371.230180 87 graph\_launch.cc:671] Fallback to op-by-op mode because memset node breaks graph update



```
[55]: y_pred_class = np.argmax(y_pred, axis=1)

y_test = np.argmax(y_test, axis=1)
```

```

accuracy = accuracy_score(y_test, y_pred_class)
print('Accuracy: %f' % accuracy)

precision = precision_score(y_test, y_pred_class, average='weighted')
print('Precision: %f' % precision)

recall = recall_score(y_test, y_pred_class, average='weighted')
print('Recall: %f' % recall)

f1 = f1_score(y_test, y_pred_class, average='weighted')
print('F1 score: %f' % f1)

```

Accuracy: 0.585861  
 Precision: 0.618950  
 Recall: 0.585861  
 F1 score: 0.582167

The same model performed better when training data is split according to driver instead of random split. A comparison between the two models is given below

	accuracy	loss	val_accuracy	val_loss	Kaggle Private Score	Kaggle Public Score
Model 1	0.4512	1.6058	0.7946	0.7764	2.24537	2.46848
Model 2	0.4960	1.4666	0.5859	1.5672	1.62022	1.71969

## 1.6 Hyperparameter Tuning

I will add a new layer with 1024 filters and use Adam optimizer to see how this model performs

```

[63]: def get_new_model():
    model = Sequential()
    # 64 conv2d filters with relu
    model.add(Conv2D(filters=64, kernel_size=2, padding='same',
    ↪activation='relu', input_shape=(64,64,3),
    ↪kernel_initializer='glorot_normal'))
    model.add(MaxPooling2D(pool_size=2))

    # 128 conv2d filters with relu
    model.add(Conv2D(filters=128, kernel_size=2, padding='same',
    ↪activation='relu', kernel_initializer='glorot_normal'))
    model.add(MaxPooling2D(pool_size=2))

    # 256 conv2d filters with relu
    model.add(Conv2D(filters=256, kernel_size=2, padding='same',
    ↪activation='relu', kernel_initializer='glorot_normal'))
    model.add(MaxPooling2D(pool_size=2))

```

```

    # 512 conv2d filters with relu
    model.add(Conv2D(filters=512, kernel_size=2, padding='same',
↪activation='relu', kernel_initializer='glorot_normal'))
    model.add(MaxPooling2D(pool_size=2))

    # 1024 conv2d filters with relu
    model.add(Conv2D(filters=1024, kernel_size=2, padding='same',
↪activation='relu', kernel_initializer='glorot_normal'))
    model.add(MaxPooling2D(pool_size=2))

    model.add(Dropout(0.5))
    model.add(Flatten())
    model.add(Dense(500, activation='relu', kernel_initializer='glorot_normal'))
    model.add(Dropout(0.5))
    model.add(Dense(10, activation='softmax',
↪kernel_initializer='glorot_normal'))
    return model

driv_selected = ['p050', 'p015', 'p022', 'p056']
## Split the data into training and validation

X_train= []
y_train = []
X_test = []
y_test = []
D_train = []
D_test = []

for features,labels,drivers in train_image:
    if drivers in driv_selected:
        X_test.append(features)
        y_test.append(labels)
        D_test.append(drivers)

    else:
        X_train.append(features)
        y_train.append(labels)
        D_train.append(drivers)

X_train = np.array(X_train).reshape(-1,img_dim,img_dim,3)
X_test = np.array(X_test).reshape(-1,img_dim,img_dim,3)
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

model = get_new_model()
model.compile(optimizer='adam', loss='categorical_crossentropy',
↪metrics=['accuracy'])

```

```

filepath = os.path.join(MODEL_PATH, "third-{epoch:02d}-{val_accuracy:.2f}.keras")

checkpoint = ModelCheckpoint(filepath, monitor='val_accuracy', verbose=1,
    ↪ save_best_only=True, mode='max')
callbacks_list = [checkpoint]

img_data_gen = ImageDataGenerator(height_shift_range=0.5, width_shift_range=0.
    ↪ 5, zoom_range=0.5, rotation_range=30)

data_generator = img_data_gen.flow(X_train, y_train, batch_size=64)

history = model.fit(data_generator, callbacks=callbacks_list, epochs=20,
    ↪ verbose=1, validation_data=(X_test, y_test), batch_size=40, shuffle=True)

```

Epoch 1/20

3/293 18s 63ms/step - accuracy:  
0.0920 - loss: 37.1295

W0000 00:00:1714042736.664572 89 graph\_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update

199/293 9s 106ms/step -  
accuracy: 0.0994 - loss: 6.1566

W0000 00:00:1714042757.530636 88 graph\_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update

292/293 0s 99ms/step -  
accuracy: 0.1044 - loss: 5.1429

W0000 00:00:1714042767.041167 89 graph\_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update

Epoch 1: val\_accuracy improved from -inf to 0.21560, saving model to  
/kaggle/working/model/self\_trained/third-01-0.22.keras

293/293 38s 109ms/step -  
accuracy: 0.1045 - loss: 5.1273 - val\_accuracy: 0.2156 - val\_loss: 2.1515

Epoch 2/20

W0000 00:00:1714042768.259210 89 graph\_launch.cc:671] Fallback to op-by-op  
mode because memset node breaks graph update

292/293 0s 89ms/step -  
accuracy: 0.1725 - loss: 2.1766

Epoch 2: val\_accuracy improved from 0.21560 to 0.31717, saving model to  
/kaggle/working/model/self\_trained/third-02-0.32.keras

293/293 27s 90ms/step -  
accuracy: 0.1726 - loss: 2.1762 - val\_accuracy: 0.3172 - val\_loss: 1.9443

Epoch 3/20  
 292/293                    0s 88ms/step -  
 accuracy: 0.2603 - loss: 1.9625  
 Epoch 3: val\_accuracy improved from 0.31717 to 0.38164, saving model to  
 /kaggle/working/model/self\_trained/third-03-0.38.keras  
 293/293                    27s 90ms/step -  
 accuracy: 0.2604 - loss: 1.9622 - val\_accuracy: 0.3816 - val\_loss: 1.6303  
 Epoch 4/20  
 292/293                    0s 87ms/step -  
 accuracy: 0.3241 - loss: 1.7817  
 Epoch 4: val\_accuracy improved from 0.38164 to 0.38218, saving model to  
 /kaggle/working/model/self\_trained/third-04-0.38.keras  
 293/293                    27s 89ms/step -  
 accuracy: 0.3242 - loss: 1.7816 - val\_accuracy: 0.3822 - val\_loss: 2.0792  
 Epoch 5/20  
 292/293                    0s 88ms/step -  
 accuracy: 0.3660 - loss: 1.6883  
 Epoch 5: val\_accuracy improved from 0.38218 to 0.38651, saving model to  
 /kaggle/working/model/self\_trained/third-05-0.39.keras  
 293/293                    27s 90ms/step -  
 accuracy: 0.3661 - loss: 1.6881 - val\_accuracy: 0.3865 - val\_loss: 2.1545  
 Epoch 6/20  
 293/293                    0s 90ms/step -  
 accuracy: 0.4054 - loss: 1.5723  
 Epoch 6: val\_accuracy did not improve from 0.38651  
 293/293                    27s 91ms/step -  
 accuracy: 0.4055 - loss: 1.5722 - val\_accuracy: 0.3692 - val\_loss: 1.9330  
 Epoch 7/20  
 292/293                    0s 88ms/step -  
 accuracy: 0.4363 - loss: 1.5113  
 Epoch 7: val\_accuracy improved from 0.38651 to 0.44664, saving model to  
 /kaggle/working/model/self\_trained/third-07-0.45.keras  
 293/293                    27s 90ms/step -  
 accuracy: 0.4364 - loss: 1.5110 - val\_accuracy: 0.4466 - val\_loss: 1.7210  
 Epoch 8/20  
 292/293                    0s 88ms/step -  
 accuracy: 0.4725 - loss: 1.4416  
 Epoch 8: val\_accuracy did not improve from 0.44664  
 293/293                    27s 89ms/step -  
 accuracy: 0.4726 - loss: 1.4413 - val\_accuracy: 0.4382 - val\_loss: 2.0683  
 Epoch 9/20  
 292/293                    0s 88ms/step -  
 accuracy: 0.5155 - loss: 1.3403  
 Epoch 9: val\_accuracy improved from 0.44664 to 0.52140, saving model to  
 /kaggle/working/model/self\_trained/third-09-0.52.keras  
 293/293                    27s 89ms/step -  
 accuracy: 0.5155 - loss: 1.3401 - val\_accuracy: 0.5214 - val\_loss: 1.4598  
 Epoch 10/20



```

293/293          0s 87ms/step -
accuracy: 0.5410 - loss: 1.2710
Epoch 10: val_accuracy did not improve from 0.52140
293/293          26s 88ms/step -
accuracy: 0.5410 - loss: 1.2709 - val_accuracy: 0.5008 - val_loss: 2.3497
Epoch 11/20
292/293          0s 88ms/step -
accuracy: 0.5817 - loss: 1.1643
Epoch 11: val_accuracy improved from 0.52140 to 0.56771, saving model to
/kaggle/working/model/self_trained/third-11-0.57.keras
293/293          27s 89ms/step -
accuracy: 0.5817 - loss: 1.1643 - val_accuracy: 0.5677 - val_loss: 1.7413
Epoch 12/20
292/293          0s 86ms/step -
accuracy: 0.6065 - loss: 1.1117
Epoch 12: val_accuracy improved from 0.56771 to 0.57042, saving model to
/kaggle/working/model/self_trained/third-12-0.57.keras
293/293          26s 88ms/step -
accuracy: 0.6066 - loss: 1.1116 - val_accuracy: 0.5704 - val_loss: 1.4773
Epoch 13/20
292/293          0s 88ms/step -
accuracy: 0.6101 - loss: 1.0923
Epoch 13: val_accuracy improved from 0.57042 to 0.61701, saving model to
/kaggle/working/model/self_trained/third-13-0.62.keras
293/293          27s 89ms/step -
accuracy: 0.6102 - loss: 1.0923 - val_accuracy: 0.6170 - val_loss: 1.4870
Epoch 14/20
292/293          0s 87ms/step -
accuracy: 0.6284 - loss: 1.0530
Epoch 14: val_accuracy improved from 0.61701 to 0.62378, saving model to
/kaggle/working/model/self_trained/third-14-0.62.keras
293/293          27s 89ms/step -
accuracy: 0.6284 - loss: 1.0529 - val_accuracy: 0.6238 - val_loss: 1.2259
Epoch 15/20
292/293          0s 90ms/step -
accuracy: 0.6365 - loss: 1.0242
Epoch 15: val_accuracy improved from 0.62378 to 0.62920, saving model to
/kaggle/working/model/self_trained/third-15-0.63.keras
293/293          27s 91ms/step -
accuracy: 0.6365 - loss: 1.0242 - val_accuracy: 0.6292 - val_loss: 1.2469
Epoch 16/20
292/293          0s 87ms/step -
accuracy: 0.6552 - loss: 0.9848
Epoch 16: val_accuracy did not improve from 0.62920
293/293          27s 88ms/step -
accuracy: 0.6552 - loss: 0.9847 - val_accuracy: 0.5756 - val_loss: 2.1654
Epoch 17/20
292/293          0s 87ms/step -

```

```

accuracy: 0.6715 - loss: 0.9381
Epoch 17: val_accuracy improved from 0.62920 to 0.65791, saving model to
/kaggle/working/model/self_trained/third-17-0.66.keras
293/293          27s 89ms/step -
accuracy: 0.6715 - loss: 0.9381 - val_accuracy: 0.6579 - val_loss: 1.2424
Epoch 18/20
292/293          0s 87ms/step -
accuracy: 0.6791 - loss: 0.9219
Epoch 18: val_accuracy did not improve from 0.65791
293/293          27s 89ms/step -
accuracy: 0.6791 - loss: 0.9220 - val_accuracy: 0.6029 - val_loss: 2.6826
Epoch 19/20
292/293          0s 88ms/step -
accuracy: 0.6745 - loss: 0.9365
Epoch 19: val_accuracy did not improve from 0.65791
293/293          27s 89ms/step -
accuracy: 0.6746 - loss: 0.9364 - val_accuracy: 0.5320 - val_loss: 2.1930
Epoch 20/20
293/293          0s 88ms/step -
accuracy: 0.6890 - loss: 0.8920
Epoch 20: val_accuracy did not improve from 0.65791
293/293          27s 89ms/step -
accuracy: 0.6890 - loss: 0.8920 - val_accuracy: 0.6227 - val_loss: 1.5915

```

```

[64]: #load checkpoint with best val_accuracy
BEST_MODEL = os.path.join(MODEL_PATH, "third-17-0.66.keras")
model = load_model(BEST_MODEL)
model.summary()
predict_and_write_results(model, test_images, files,
    ↪ 'kaggle_self_trained_third_model.csv')

```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
conv2d_22 (Conv2D)	(None, 64, 64, 64)	832
max_pooling2d_22 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_23 (Conv2D)	(None, 32, 32, 128)	32,896
max_pooling2d_23 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_24 (Conv2D)	(None, 16, 16, 256)	131,328
max_pooling2d_24 (MaxPooling2D)	(None, 8, 8, 256)	0

conv2d_25 (Conv2D)	(None, 8, 8, 512)	524,800
max_pooling2d_25 (MaxPooling2D)	(None, 4, 4, 512)	0
conv2d_26 (Conv2D)	(None, 4, 4, 1024)	2,098,176
max_pooling2d_26 (MaxPooling2D)	(None, 2, 2, 1024)	0
dropout_10 (Dropout)	(None, 2, 2, 1024)	0
flatten_5 (Flatten)	(None, 4096)	0
dense_10 (Dense)	(None, 500)	2,048,500
dropout_11 (Dropout)	(None, 500)	0
dense_11 (Dense)	(None, 10)	5,010

Total params: 14,524,628 (55.41 MB)

Trainable params: 4,841,542 (18.47 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 9,683,086 (36.94 MB)

54/2492 7s 3ms/step




W0000 00:00:1714043305.647168 88 graph\_launch.cc:671] Fallback to op-by-op mode because memset node breaks graph update

2492/2492 9s 3ms/step

W0000 00:00:1714043314.027145 89 graph\_launch.cc:671] Fallback to op-by-op mode because memset node breaks graph update

/tmp/ipykernel\_34/950582733.py:18: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```
df = pd.concat([df, pd.DataFrame([new_row])], ignore_index=True)
```

Submission and Description	Private Score 	Public Score 	Selected
 <b>kaggle_self_trained_third_model.csv</b> <small>Complete (after deadline) · now</small>	1.10784	1.30779	<input type="checkbox"/>

This model performed better than the last 2 models after the hyper parameter tuning

	accuracy	loss	val_accuracy	val_loss	Kaggle Private Score	Kaggle Public Score
Model 1	0.4512	1.6058	0.7946	0.7764	2.24537	2.46848
Model 2	0.4960	1.4666	0.5859	1.5672	1.62022	1.71969
Model 3	0.6715	0.9381	0.6579	1.2424	1.10784	1.30779

## 1.7 Results Comparision

Apart from the self trained models, I have tried transfer learning using Mobilenet and Resnet models. I am not adding those details in this notebook file since it will increase the size of this notebook.

The corresponding Jupyter notebooks are available on github:  
[https://github.com/krishnakuruvadi/5511\\_Final](https://github.com/krishnakuruvadi/5511_Final)

Comparision between the best models is given below

	accuracy	loss	val_accuracy	val_loss	Kaggle Private Score	Kaggle Public Score
Self trained	0.6715	0.9381	0.6579	1.2424	1.10784	1.30779
Mobilenet	0.8816	0.3670	0.8621	0.4850	0.41831	0.60318
Resnet	0.7843	0.6402	0.8602	0.4452	0.41789	0.52613

## 1.8 Findings From Results

Mobilenet and Resnet achieved almost same results and the results are better then the self trained model. When we compare the self trained models with these pre-trained models, we see the number of layers and parameters are more. Also the weights in the pre-trained models are set based on a lot of training data.

## 1.9 Conclusion

- Pre-trained models work well for such usecase when compared to a general self trained model
- We could take the pre-trained models and add more layers or use the pre-trained models as reference to come up with a better self trained model
- Training takes more time and without GPUs working on this kind of tasks is very time consuming

What I would like to try:








- I would try to use other recent image processing pre-trained models like YOLO
- Fine tune the hyper parameters further
- Change image size/resolution to 224 instead of 64

## 1.10 Deliverables


Github: [https://github.com/krishnakuruvadi/5511\\_Final](https://github.com/krishnakuruvadi/5511_Final)

Presentation: [https://github.com/krishnakuruvadi/5511\\_Final](https://github.com/krishnakuruvadi/5511_Final)

Kaggle scores of all models tried (including the ones not available on github):

Submission and Description	Private Score ⓘ	Public Score ⓘ	Selected
 <b>kaggle_self_trained_second_model.csv</b> Complete (after deadline) · 27m ago	<b>1.62022</b>	<b>1.71969</b>	<input type="checkbox"/>
 <b>kaggle_self_trained_first_model.csv</b> Complete (after deadline) · 1h ago	<b>2.24537</b>	<b>2.46848</b>	<input type="checkbox"/>
 <b>self_trained_test_result_4.csv</b> Complete (after deadline) · 12h ago	<b>1.72830</b>	<b>1.75492</b>	<input type="checkbox"/>
 <b>self_trained_test_result_3.csv</b> Complete (after deadline) · 18h ago	<b>2.01536</b>	<b>2.81197</b>	<input type="checkbox"/>
 <b>resnet_sgd_1.csv</b> Complete (after deadline) · 1d ago	<b>0.41789</b>	<b>0.52613</b>	<input type="checkbox"/>
 <b>mobilenet_sgd_1.csv</b> Complete (after deadline) · 3d ago	<b>0.41831</b>	<b>0.60318</b>	<input type="checkbox"/>
 <b>self_trained_test_result_2.csv</b> Complete (after deadline) · 4d ago	<b>1.94104</b>	<b>1.38854</b>	<input type="checkbox"/>

Probable leaderboard place based on best model results (Resnet) is ~271/1439:

270	↗ 21	FromChn		0.41577	2	8y
271	↗ 19	TeamD	 	0.41922	21	8y

## 1.11 References

- RESNET: [https://en.wikipedia.org/wiki/Residual\\_neural\\_network](https://en.wikipedia.org/wiki/Residual_neural_network)
- Mobilenet: <https://keras.io/api/applications/mobilenet/>
- State Farm Kaggle competition: <https://www.kaggle.com/competitions/state-farm-distracted-driver-detection>